



```
In [1]: import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import sqlite3
import csv
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from wordcloud import WordCloud
import re
import os
from sqlalchemy import create_engine # database connection
import datetime as dt
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem.snowball import SnowballStemmer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.multiclass import OneVsRestClassifier
from sklearn.linear_model import SGDClassifier
from sklearn import metrics
from sklearn.metrics import f1_score, precision_score, recall_score
from sklearn import svm
from sklearn.linear_model import LogisticRegression
from skmultilearn.adapt import mlknn
from skmultilearn.problem_transform import ClassifierChain
from skmultilearn.problem_transform import BinaryRelevance
from skmultilearn.problem_transform import LabelPowerset
from sklearn.naive_bayes import GaussianNB
from datetime import datetime
```

## Stack Overflow: Tag Prediction

# 1. Business Problem

## 1.1 Description

### Description

Stack Overflow is the largest, most trusted online community for developers to learn, share their programming knowledge, and build their careers.

Stack Overflow is something which every programmer use one way or another. Each month, over 50 million developers come to Stack Overflow to learn, share their knowledge, and build their careers. It features questions and answers on a wide range of topics in computer programming. The website serves as a platform for users to ask and answer questions, and, through membership and active participation, to vote questions and answers up or down and edit questions and answers in a fashion similar to a wiki or Digg. As of April 2014 Stack Overflow has over 4,000,000 registered users, and it exceeded 10,000,000 questions in late August 2015. Based on the type of tags assigned to questions, the top eight most discussed topics on the site are: Java, JavaScript, C#, PHP, Android, jQuery, Python and HTML.

### Problem Statement

Suggest the tags based on the content that was there in the question posted on Stackoverflow.

**Source:** <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/>

## 1.2 Source / useful links

Data Source : <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>

Youtube : <https://youtu.be/nNDqbUhtIRg>

Research paper : <https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/tagging-1.pdf>

Research paper : <https://dl.acm.org/citation.cfm?id=2660970&dl=ACM&coll=DL>

## 1.3 Real World / Business Objectives and Constraints

1. Predict as many tags as possible with high precision and recall.
2. Incorrect tags could impact customer experience on StackOverflow.
3. No strict latency constraints.

## 2. Machine Learning problem

### 2.1 Data

#### 2.1.1 Data Overview

Refer: <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>

All of the data is in 2 files: Train and Test.

**Train.csv** contains 4 columns: Id, Title, Body, Tags.

**Test.csv** contains the same columns but without the Tags, which you are to predict.

**Size of Train.csv** - 6.75GB

**Size of Test.csv** - 2GB

**Number of rows in Train.csv** = 6034195

The questions are randomized and contains a mix of verbose text sites as well as sites related to math and programming. The number of questions from each site may vary, and no filtering has been performed on the questions (such as closed questions).

### **Data Field Explanation**

Dataset contains 6,034,195 rows. The columns in the table are:

**Id** - Unique identifier for each question

**Title** - The question's title

**Body** - The body of the question

**Tags** - The tags associated with the question in a space-separated format (all lowercase, should not contain tabs '\t' or ampersands '&')

### **2.1.2 Example Data point**

**Title:** Implementing Boundary Value Analysis of Software Testing in a C++ program?

**Body :**

```

#include<
iostream>\n
#include<
stdlib.h>\n\n
using namespace std;\n\n
int main()\n
{\n
    int n,a[n],x,c,u[n],m[n],e[n][4];\n

    cout<<"Enter the number of variables";\n
    cin>>n;\n\n
    cout<<"Enter the Lower, and Upper Limits
of the variables";\n
    for(int y=1; y<n+1; y++)\n
    {\n
        cin>>m[y];\n
        cin>>u[y];\n
    }\n
    for(x=1; x<n+1; x++)\n
    {\n
        a[x] = (m[x] + u[x])/2;\n
    }\n
    c=(n*4)-4;\n
    for(int a1=1; a1<n+1; a1++)\n
    {\n\n
        e[a1][0] = m[a1];\n
        e[a1][1] = m[a1]+1;\n
        e[a1][2] = u[a1]-1;\n
        e[a1][3] = u[a1];\n
    }\n
    for(int i=1; i<n+1; i++)\n
    {\n

```

```

        for(int l=1; l<=i; l++)\n
        {\n
            if(l!=1)\n
            {\n
                cout<<a[l]<<"\\t";\n

            }\n
        }\n
        for(int j=0; j<4; j++)\n
        {\n
            cout<<e[i][j];\n
            for(int k=0; k<n-(i+1); k++)\n

            {\n
                cout<<a[k]<<"\\t";\n

            }\n
            cout<<"\\n";\n
        }\n
    }\n\n
    system("PAUSE");\n
    return 0;    \n
}\n

```

\n\n

The answer should come in the form of a table like  
 \n\n

1

50

50\n

2	50	50\n
99	50	50\n
100	50	50\n
50	1	50\n
50	2	50\n
50	99	50\n
50	100	50\n
50	50	1\n
50	50	2\n
50	50	99\n
50	50	100\n

\n\n

if the no of inputs is 3 and their ranges are\n

1,100\n

1,100\n

1,100\n

(could be varied too)

\n\n

The output is not coming,can anyone correct the code or tell me

what\'s wrong?

\n'

**Tags** : 'c++ c'

## 2.2 Mapping the real-world problem to a Machine Learning Problem

## 2.2.1 Type of Machine Learning Problem

It is a multi-label classification problem

**Multi-label Classification:** Multilabel classification assigns to each sample a set of target labels. This can be thought as predicting properties of a data-point that are not mutually exclusive, such as topics that are relevant for a document. A question on Stackoverflow might be about any of C, Pointers, FileIO and/or memory-management at the same time or none of these.

\_\_Credit\_\_: <http://scikit-learn.org/stable/modules/multiclass.html>

## 2.2.2 Performance metric

**Micro-Averaged F1-Score (Mean F Score)** : The F1 score can be interpreted as a weighted average of the precision and recall, where an F1 score reaches its best value at 1 and worst score at 0. The relative contribution of precision and recall to the F1 score are equal. The formula for the F1 score is:

$$F1 = 2 * (precision * recall) / (precision + recall)$$

In the multi-class and multi-label case, this is the weighted average of the F1 score of each class.

**'Micro f1 score':**

Calculate metrics globally by counting the total true positives, false negatives and false positives. This is a better metric when we have class imbalance.

**'Macro f1 score':**

Calculate metrics for each label, and find their unweighted mean. This does not take label imbalance into account.

<https://www.kaggle.com/wiki/MeanFScore>

[http://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1\\_score.html](http://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1_score.html)

**Hamming loss** : The Hamming loss is the fraction of labels that are incorrectly predicted.

<https://www.kaggle.com/wiki/HammingLoss>



## 3. Exploratory Data Analysis

### 3.1 Data Loading and Cleaning

#### 3.1.1 Using Pandas with SQLite to Load the data

```
In [2]: #Creating db file from csv
#Learn SQL: https://www.w3schools.com/sql/default.asp
if not os.path.isfile('train.db'):
    start = datetime.now()
    disk_engine = create_engine('sqlite:///train.db')
    start = dt.datetime.now()
    chunksize = 180000
    j = 0
    index_start = 1
    for df in pd.read_csv('Train.csv', names=['Id', 'Title', 'Body', 'Tags'], chunksize=chunksize, iterator=True, encoding='utf-8', ):
        df.index += index_start
        j+=1
        print('{} rows'.format(j*chunksize))
        df.to_sql('data', disk_engine, if_exists='append')
        index_start = df.index[-1] + 1
    print("Time taken to run this cell :", datetime.now() - start)
```

#### 3.1.2 Counting the number of rows

```
In [3]: if os.path.isfile('train.db'):
    start = datetime.now()
    con = sqlite3.connect('train.db')
    num_rows = pd.read_sql_query("""SELECT count(*) FROM data""", con)
    #Always remember to close the database
```

```

        print("Number of rows in the database :", "\n", num_rows['count(*)'].
values[0])
        con.close()
        print("Time taken to count the number of rows :", datetime.now() -
start)
    else:
        print("Please download the train.db file from drive or run the abov
e cell to generate train.db file")

```

Number of rows in the database :

6034196

Time taken to count the number of rows : 0:00:00.337998

### 3.1.3 Checking for duplicates

```

In [4]: #Learn SQL: https://www.w3schools.com/sql/default.asp
if os.path.isfile('train.db'):
    start = datetime.now()
    con = sqlite3.connect('train.db')
    df_no_dup = pd.read_sql_query('SELECT Title, Body, Tags, COUNT(*) a
s cnt_dup FROM data WHERE Tags IS NOT NULL \
                                GROUP BY Title, Body, Tags LIMIT 5000
00', con)
    con.close()
    print("Time taken to run this cell :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the firs
t to generate train.db file")

```

Time taken to run this cell : 0:01:17.795018

```

In [5]: df_no_dup.head()
# we can observe that there are duplicates

```

Out[5]:

Title	Body	Tags	cnt_dup
-------	------	------	---------

	Title	Body	Tags	cnt_dup
0	Implementing Boundary Value Analysis of S...	<pre>&lt;code&gt;#include&lt;iosstream&gt;\n#include&lt;...&lt;/pre&gt;</pre>	c++ c	1
1	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamicall...</p>	c# silverlight data-binding	1
2	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamicall...</p>	c# silverlight data-binding columns	1
3	java.lang.NoClassDefFoundError: javax/serv...	<p>I followed the guide in &lt;a href="http://sta...</p>	jsp jstl	1
4	java.sql.SQLException:[Microsoft][ODBC Dri...	<p>I use the following code&lt;/p&gt;\n\n<pre>&lt;code&gt;...&lt;/code&gt;...</pre></p>	java jdbc	2

```
In [6]: print("number of duplicate questions :", num_rows['count(*)'].values[0]
- df_no_dup.shape[0], "(", (1-((df_no_dup.shape[0])/(num_rows['count(*)'
].values[0]))) * 100, "% )")
```

number of duplicate questions : 5534196 ( 91.71389195843159 % )

```
In [7]: # number of times each question appeared in our database
df_no_dup.cnt_dup.value_counts()
```

```
Out[7]: 1    315993
2    150976
3     33020
4         9
6         1
5         1
Name: cnt_dup, dtype: int64
```

```
In [8]: start = datetime.now()
df_no_dup["tag_count"] = df_no_dup["Tags"].apply(lambda text: len(text.
split(" ")))
```

```
# adding a new feature number of tags per question
print("Time taken to run this cell :", datetime.now() - start)
df_no_dup.head()
```

Time taken to run this cell : 0:00:00.376469

Out[8]:

	Title	Body	Tags	cnt_dup	tag
0	Implementing Boundary Value Analysis of S...	<pre>#include<iosstream>\n#include<...</pre>	c++ c	1	
1	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamicall...	c# silverlight data-binding	1	
2	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamicall...	c# silverlight data-binding columns	1	
3	java.lang.NoClassDefFoundError: javax/serv...	<p>I followed the guide in <a href="http://sta...	jsp jstl	1	
4	java.sql.SQLException:[Microsoft][ODBC Dri...	<p>I use the following code</p>\n\n<pre>#include<iosstream>\n#include<...</pre>	java jdbc	2	

```
In [9]: # distribution of number of tags per question
df_no_dup.tag_count.value_counts()
```

```
Out[9]: 3    144976
        2    128685
        4     99996
        1     63344
        5     62999
        Name: tag_count, dtype: int64
```

```
In [10]: #Creating a new database with no duplicates
if not os.path.isfile('train_no_dup.db'):
    disk_dup = create_engine("sqlite:///train_no_dup.db")
```

```
no_dup = pd.DataFrame(df_no_dup, columns=['Title', 'Body', 'Tags'])
no_dup.to_sql('no_dup_train', disk_dup)
```

```
In [11]: #This method seems more appropriate to work with this much data.
#creating the connection with database file.
if os.path.isfile('train_no_dup.db'):
    start = datetime.now()
    con = sqlite3.connect('train_no_dup.db')
    tag_data = pd.read_sql_query("""SELECT Tags FROM no_dup_train WHERE
Tags IS NOT NULL""", con)
    #Always remember to close the database
    con.close()

    # Let's now drop unwanted column.
    tag_data.drop(tag_data.index[0], inplace=True)
    #Printing first 5 columns from our data frame
    tag_data.head()
    print("Time taken to run this cell :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the above
cells to generate train.db file")
```

Time taken to run this cell : 0:00:20.241672

## 3.2 Analysis of Tags

### 3.2.1 Total number of unique tags

```
In [12]: # Importing & Initializing the "CountVectorizer" object, which
#is scikit-learn's bag of words tool.

#by default 'split()' will tokenize each tag using space.
vectorizer = CountVectorizer(tokenizer = lambda x: x.split(), max_features=20000,
ngram_range=(1,2))
# fit_transform() does two functions: First, it fits the model
# and learns the vocabulary; second, it transforms our training data
```

```
# into feature vectors. The input to fit_transform should be a list of strings.
tag_dtm = vectorizer.fit_transform(tag_data['Tags'])
```

```
In [13]: print("Number of data points :", tag_dtm.shape[0])
         print("Number of unique tags :", tag_dtm.shape[1])
```

```
Number of data points : 4206307
Number of unique tags : 20000
```

```
In [14]: #'get_feature_name()' gives us the vocabulary.
         tags = vectorizer.get_feature_names()
         #Lets look at the tags we have.
         print("Some of the tags we have :", tags[:10])
```

```
Some of the tags we have : ['.bash-profile', '.each', '.htaccess', '.ht
access apache2', '.htaccess cakephp', '.htaccess codeigniter', '.htacce
ss magento', '.htaccess mod-rewrite', '.htaccess redirect', '.htaccess
rewrite']
```

### 3.2.3 Number of times a tag appeared

```
In [15]: # https://stackoverflow.com/questions/15115765/how-to-access-sparse-mat
         rix-elements
         #Lets now store the document term matrix in a dictionary.
         freqs = tag_dtm.sum(axis=0).A1
         result = dict(zip(tags, freqs))
```


```
In [16]: #Saving this dictionary to csv files.
         if not os.path.isfile('tag_counts_dict_dtm.csv'):
             with open('tag_counts_dict_dtm.csv', 'w') as csv_file:
                 writer = csv.writer(csv_file)
                 for key, value in result.items():
                     writer.writerow([key, value])
         tag_df = pd.read_csv("tag_counts_dict_dtm.csv", names=['Tags', 'Counts'
])
         tag_df.head()
```

Out[16]:


	Tags	Counts
0	storyboard	1939
1	html5 html5-video	273
2	html-tag	79
3	php opencart	128
4	android css	266

```
In [17]: tag_df_sorted = tag_df.sort_values(['Counts'], ascending=False)
tag_counts = tag_df_sorted['Counts'].values
```

```
In [18]: plt.plot(tag_counts)
plt.title("Distribution of number of times tag appeared questions")
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
```



```
In [19]: plt.plot(tag_counts[0:10000])
plt.title('first 10k tags: Distribution of number of times tag appeared
questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:10000:25]), tag_counts[0:10000:25])
```




```
400 [331505  46769  25375  20180  15380  13157  11023  10029   9159    8
197
7528   6704   6207   5817   5421   5128   4890   4526   4314   4201
4050   3874   3728   3609   3483   3375   3235   3091   2986   2900
2756   2698   2617   2531   2449   2399   2331   2266   2211   2162
2107   2061   2013   1965   1929   1876   1839   1804   1757   1723
1695   1662   1631   1601   1555   1525   1491   1464   1444   1419
1390   1363   1339   1317   1300   1279   1255   1240   1225   1210
1191   1178   1163   1153   1137   1121   1109   1090   1077   1065
1051   1038   1029   1016   1005   985    974   962   951   941
928    918    908    894    886    878    871    862    847    839
832    822    814    806    798    788    781    776    768    759
747    742    735    731    725    717    711    702    694    689
685    680    673    666    659    654    649    644    638    633
---
```



629	623	617	613	607	603	597	590	585	582
578	573	568	562	558	553	549	544	541	538
532	528	524	521	516	513	510	506	503	499
495	492	489	485	482	479	475	473	468	466
462	459	456	453	449	447	444	442	439	437
433	431	429	426	424	422	419	416	414	412
409	406	403	401	399	396	393	390	388	387
385	383	380	378	375	374	372	370	369	366
365	363	361	359	356	354	352	351	349	348
346	344	342	341	340	338	336	335	333	331
330	328	326	324	323	321	319	318	316	314
313	311	310	308	307	306	304	303	301	300
299	298	296	295	294	293	292	291	289	288
286	285	284	283	282	280	279	278	277	276
275	273	272	271	270	269	268	267	265	264
263	262	260	259	259	258	257	256	255	253
252	251	251	250	249	248	248	247	246	245
244	243	242	241	240	239	238	238	237	236
234	234	233	232	231	231	230	229	228	227
227	226	225	224	223	222	221	220	220	219
218	217	217	216	215	214	214	213	213	212
211	210	209	209	208	207	206	206	205	205
204	203	202	201	201	200	199	199	198	197
197	196	196	195	194	194	193	193	192	192
191	190	190	189	188	188	188	187	186	186
185	184	184	183	183	182	182	181	180	180
179	179	178	178	177	177	176	176	175	175]

```
In [20]: plt.plot(tag_counts[0:1000])
plt.title('first 1k tags: Distribution of number of times tag appeared
questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:1000:5]), tag_counts[0:1000:5])
```




```

200 [331505 221533 122769 95160 69168 46769 42006 37868 33373 31
069
25375 23585 22747 21942 21006 20180 19161 18432 17640 15671
15380 15087 14108 13532 13364 13157 12407 11897 11454 11208
11023 10741 10525 10288 10205 10029 9981 9787 9523 9265
9159 9079 8833 8698 8545 8197 8065 7933 7828 7684
7528 7265 7115 7044 6856 6704 6628 6513 6447 6291
6207 6139 6039 5971 5866 5817 5757 5669 5574 5498
5421 5399 5360 5283 5209 5128 5083 5020 4988 4939
4890 4785 4663 4593 4548 4526 4501 4447 4392 4335
4314 4289 4273 4241 4233 4201 4184 4157 4142 4088
4050 4003 3971 3947 3919 3874 3849 3822 3799 3757
3728 3703 3685 3660 3636 3609 3583 3563 3521 3509
3483 3468 3447 3422 3402 3375 3333 3309 3293 3272
3235 3198 3174 3148 3108 3091 3073 3050 3032 3008
2986 2983 2953 2935 2920 2900 2878 2843 2819 2784
2756 2749 2733 2724 2716 2698 2679 2667 2652 2627
2617 2604 2598 2586 2548 2531 2520 2502 2489 2470
2449 2441 2431 2416 2408 2399 2387 2371 2363 2346
2331 2312 2301 2295 2284 2266 2259 2248 2242 2222
2211 2206 2197 2186 2174 2162 2146 2138 2127 2117]

```

```
In [21]: plt.plot(tag_counts[0:500])
plt.title('first 500 tags: Distribution of number of times tag appeared
questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:500:5]), tag_counts[0:500:5])
```




```
100 [331505 221533 122769 95160 69168 46769 42006 37868 33373 31
069
25375 23585 22747 21942 21006 20180 19161 18432 17640 15671
15380 15087 14108 13532 13364 13157 12407 11897 11454 11208
11023 10741 10525 10288 10205 10029 9981 9787 9523 9265
9159 9079 8833 8698 8545 8197 8065 7933 7828 7684
7528 7265 7115 7044 6856 6704 6628 6513 6447 6291
6207 6139 6039 5971 5866 5817 5757 5669 5574 5498
5421 5399 5360 5283 5209 5128 5083 5020 4988 4939
4890 4785 4663 4593 4548 4526 4501 4447 4392 4335
4314 4289 4273 4241 4233 4201 4184 4157 4142 4088]
```

```
In [22]: plt.plot(tag_counts[0:100], c='b')
plt.scatter(x=list(range(0,100,5)), y=tag_counts[0:100:5], c='orange',
label="quantiles with 0.05 intervals")
# quantiles with 0.25 difference
plt.scatter(x=list(range(0,100,25)), y=tag_counts[0:100:25], c='m', lab
el = "quantiles with 0.25 intervals")

for x,y in zip(list(range(0,100,25)), tag_counts[0:100:25]):
    plt.annotate(s="({} , {})".format(x,y), xy=(x,y), xytext=(x-0.05, y
+500))

plt.title('first 100 tags: Distribution of number of times tag appeared
questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.legend()
plt.show()
print(len(tag_counts[0:100:5]), tag_counts[0:100:5])
```



```
20 [331505 221533 122769 95160 69168 46769 42006 37868 33373 310
69
25375 23585 22747 21942 21006 20180 19161 18432 17640 15671]
```

```
In [23]: # Store tags greater than 10K in one list
lst_tags_gt_10k = tag_df[tag_df.Counts>10000].Tags
#Print the length of the list
print ('{} Tags are used more than 10000 times'.format(len(lst_tags_gt_10k)))
# Store tags greater than 100K in one list
lst_tags_gt_100k = tag_df[tag_df.Counts>100000].Tags
#Print the length of the list.
print ('{} Tags are used more than 100000 times'.format(len(lst_tags_gt_100k)))
```

```
178 Tags are used more than 10000 times
14 Tags are used more than 100000 times
```

#### Observations:

1. There are total 178 tags which are used more than 10000 times.
2. 14 tags are used more than 100000 times.
3. Most frequent tag (i.e. c#) is used 331505 times.
4. Since some tags occur much more frequently than others, Micro-averaged F1-score is the appropriate metric for this problem.

### 3.2.4 Tags Per Question

```
In [24]: #Storing the count of tag in each question in list 'tag_count'
tag_quest_count = tag_dtm.sum(axis=1).tolist()
#Converting list of lists into single list, we will get [[3], [4], [2], [2], [3]] and we are converting this to [3, 4, 2, 2, 3]
tag_quest_count=[int(j) for i in tag_quest_count for j in i]
print ('We have total {} datapoints.'.format(len(tag_quest_count)))


print(tag_quest_count[:5])
```

```
We have total 4206307 datapoints.
[5, 6, 3, 3, 4]
```

```
In [25]: print( "Maximum number of tags per question: %d"%max(tag_quest_count))
print( "Minimum number of tags per question: %d"%min(tag_quest_count))
print( "Avg. number of tags per question: %f"% ((sum(tag_quest_count)*
1.0)/len(tag_quest_count)))
```

Maximum number of tags per question: 9  
Minimum number of tags per question: 0  
Avg. number of tags per question: 3.669640

```
In [26]: sns.countplot(tag_quest_count, palette='gist_rainbow')
plt.title("Number of tags in the questions ")
plt.xlabel("Number of Tags")
plt.ylabel("Number of questions")
plt.show()
```



#### Observations:

1. Maximum number of tags per question: 9
2. Minimum number of tags per question: 0
3. Avg. number of tags per question: 3.669


4. Most of the questions are having 3 or 5 tags.

### 3.2.5 Most Frequent Tags

```
In [27]: # Plotting word cloud
start = datetime.now()

# Lets first convert the 'result' dictionary to 'list of tuples'
tup = dict(result.items())
#Initializing WordCloud using frequencies of tags.
wordcloud = WordCloud(    background_color='black',
                          width=1600,
                          height=800,
                          ).generate_from_frequencies(tup)

fig = plt.figure(figsize=(30,20))
plt.imshow(wordcloud)
plt.axis('off')
plt.tight_layout(pad=0)
fig.savefig("tag.png")
plt.show()
print("Time taken to run this cell :", datetime.now() - start)
```



Time taken to run this cell : 0:00:05.020815


### Observations:

A look at the word cloud shows that "c#", "java", "php", "asp.net", "javascript", "c++" are some of the most frequent tags.

### 3.2.6 The top 20 tags

```
In [28]: i=np.arange(30)
tag_df_sorted.head(30).plot(kind='bar')
plt.title('Frequency of top 20 tags')
plt.xticks(i, tag_df_sorted['Tags'])
plt.xlabel('Tags')
plt.ylabel('Counts')
plt.show()
```





#### Observations:

1. Majority of the most frequent tags are programming language.
2. C# is the top most frequent programming language.
3. Android, IOS, Linux and windows are among the top most frequent operating systems.

### 3.3 Cleaning and preprocessing of Questions

#### 3.3.1 Preprocessing

1. Sample 0.5M data points
2. Separate out code-snippets from Body

3. Remove Special characters from Question title and description (not in code)
4. Remove stop words (Except 'C')
5. Remove HTML Tags
6. Convert all the characters into small letters
7. Use SnowballStemmer to stem the words

```
In [29]: import nltk
nltk.download('stopwords')

def striphtml(data):
    cleanr = re.compile('<.*?>')
    cleantext = re.sub(cleanr, ' ', str(data))
    return cleantext
stop_words = set(stopwords.words('english'))
stemmer = SnowballStemmer("english")

[nltk_data] Downloading package stopwords to
[nltk_data] /home/sujit_biswal624882/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
In [30]: #http://www.sqlitetutorial.net/sqlite-python/create-tables/
def create_connection(db_file):
    """ create a database connection to the SQLite database
        specified by db_file
    :param db_file: database file
    :return: Connection object or None
    """
    try:
        conn = sqlite3.connect(db_file)
        return conn
    except Error as e:
        print(e)

    return None

def create_table(conn, create_table_sql):
    """ create a table from the create_table_sql statement
```

```

:param conn: Connection object
:param create_table_sql: a CREATE TABLE statement
:return:
"""
try:
    c = conn.cursor()
    c.execute(create_table_sql)
except Error as e:
    print(e)

def checkTableExists(dbcon):
    cursr = dbcon.cursor()
    str = "select name from sqlite_master where type='table'"
    table_names = cursr.execute(str)
    print("Tables in the database:")
    tables = table_names.fetchall()
    print(tables[0][0])
    return(len(tables))

def create_database_table(database, query):
    conn = create_connection(database)
    if conn is not None:
        create_table(conn, query)
        checkTableExists(conn)
    else:
        print("Error! cannot create the database connection.")
    conn.close()

sql_create_table = """CREATE TABLE IF NOT EXISTS QuestionsProcessed (question text NOT NULL, code text, tags text, words_pre integer, words_post integer, is_code integer);"""
create_database_table("Processed.db", sql_create_table)

```

Tables in the database:  
QuestionsProcessed

In [31]: [# http://www.sqlitetutorial.net/sqlite-delete/](http://www.sqlitetutorial.net/sqlite-delete/)  
[# https://stackoverflow.com/questions/2279706/select-random-row-from-a-sqlite-table](https://stackoverflow.com/questions/2279706/select-random-row-from-a-sqlite-table)

```

start = datetime.now()
read_db = 'train_no_dup.db'
write_db = 'Processed.db'
if os.path.isfile(read_db):
    conn_r = create_connection(read_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        reader.execute("SELECT Title, Body, Tags From no_dup_train ORDE
R BY RANDOM() LIMIT 50000;")

if os.path.isfile(write_db):
    conn_w = create_connection(write_db)
    if conn_w is not None:
        tables = checkTableExists(conn_w)
        writer = conn_w.cursor()
        if tables != 0:
            writer.execute("DELETE FROM QuestionsProcessed WHERE 1")
            print("Cleared All the rows")
print("Time taken to run this cell :", datetime.now() - start)

```

Tables in the database:  
 QuestionsProcessed  
 Cleared All the rows  
 Time taken to run this cell : 0:00:26.870566

**we create a new data base to store the sampled and preprocessed questions**

In [32]: [#### http://www.bernzilla.com/2008/05/13/selecting-a-random-row-from-an-sqlite-table/](http://www.bernzilla.com/2008/05/13/selecting-a-random-row-from-an-sqlite-table/)

```

import nltk
nltk.download('punkt')

start = datetime.now()
preprocessed_data_list=[]
reader.fetchone()
questions_with_code=0
len_pre=0

```

```

len_post=0
questions_proccesed = 0
for row in reader:

    is_code = 0

    title, question, tags = row[0], row[1], row[2]

    if '<code>' in question:
        questions_with_code+=1
        is_code = 1
    x = len(question)+len(title)
    len_pre+=x

    code = str(re.findall(r'<code>(.*?)</code>', question, flags=re.DOTALL))

    question=re.sub('<code>(.*?)</code>', '', question, flags=re.MULTILINE|re.DOTALL)
    question=stripthtml(question.encode('utf-8'))

    title=title.encode('utf-8')

    question=str(title)+" "+str(question)
    question=re.sub(r'^[A-Za-z]+', ' ', question)
    words=word_tokenize(str(question.lower()))

    #Removing all single letter and and stopwords from question except for the letter 'c'
    question=' '.join(str(stemmer.stem(j)) for j in words if j not in stopwords and (len(j)!=1 or j=='c'))

    len_post+=len(question)
    tup = (question,code,tags,x,len(question),is_code)
    questions_proccesed += 1
    writer.execute("insert into QuestionsProcessed(question,code,tags,words_pre,words_post,is_code) values (?,?,?,?,?,?)",tup)
    if (questions_proccesed%100000==0):
        print("number of questions completed=",questions_proccesed)

```

```

no_dup_avg_len_pre=(len_pre*1.0)/questions_proccesed
no_dup_avg_len_post=(len_post*1.0)/questions_proccesed

print( "Avg. length of questions(Title+Body) before processing: %d"%no_
dup_avg_len_pre)
print( "Avg. length of questions(Title+Body) after processing: %d"%no_d
up_avg_len_post)
print( "Percent of questions containing code: %d"%((questions_with_code
*100.0)/questions_proccesed))

print("Time taken to run this cell :", datetime.now() - start)

```

```

[nltk_data] Downloading package punkt to
[nltk_data]      /home/sujit_biswal624882/nltk_data...
[nltk_data]   Package punkt is already up-to-date!

```

```

Avg. length of questions(Title+Body) before processing: 1171
Avg. length of questions(Title+Body) after processing: 325
Percent of questions containing code: 57
Time taken to run this cell : 0:01:14.824407

```

```

In [33]: # dont forget to close the connections, or else you will end up with lo
cks
conn_r.commit()
conn_w.commit()
conn_r.close()
conn_w.close()

```

```

In [34]: if os.path.isfile(write_db):
        conn_r = create_connection(write_db)
        if conn_r is not None:
            reader =conn_r.cursor()
            reader.execute("SELECT question From QuestionsProcessed LIMIT 1
0")

            print("Questions after preprocessed")
            print('='*100)
            reader.fetchone()
            for row in reader:

```

```

        print(row)
        print('-'*100)
conn_r.commit()
conn_r.close()

```

Questions after preprocessed

=====

('self sign certif work behind apach revers proxi want use apach revers  
proxi collect app server plan se ca sign ssl certif apach instanc want  
use self sign certif app server instanc apach app sever connect also en  
crypt dont want instal ca sign ssl certif app server instanc dont apach  
allow configur self sign certif app server instanc',)

-----

('python return max number array error python script connect websit via  
ftp list current version number program locat websit creat array hold v  
ersion number till script would pick largest number array tell exampl a  
rray would usual look like use return largest valu array unfortun recei  
v largest number fact latest version number nthe thing could think cut  
list array receiv past back get full version number sinc websit contain  
mani version know effici help would great appreci',)

-----

('auto start applic user login mac java program creat bundl jar file cr  
eat packag want auto start applic user login time rather manual set log  
in item go user account system pref went link make app open login use l  
aunchservic lssharedfilelist link code packag maker add xcode code post  
instal script',)

-----

('xcode ctrl drag edit code snippet xcode ctrl drag ib implement get th  
ing like great howev chang code snippet action ni want brace line aka l  
ook like',)

-----

('close select dropdown list programat javascript jquery dropdown initi  
one singl valu user click singl element remov new element ad say load a  
jax call issu server return new valu ad control problem control remain  
open updat would like close exampl http jsfiddl net vtortola cgubk exam

```
pl ajax get data probabl someth wrong call jsfiddl api show select rema
in open updat want know close dropdown list programat focus anoth input
cheer',)
```

```
-----
('mutual friend request return wrong mutual friend app show user mutual
facebook friend show friend mutual friend well get show mutual friend w
ant see friend mutual friend random peopl',)
```

```
-----
('integr jetbrain teamciti atlassian stash stash come new rest api allo
w tell stash build relat specif changeset let stash know build teamcit
i',)
```

```
-----
('use urllib submit soap request soap request known work use tool like
say soapui tri get work use urllib tri far work abl find spec document
post soap server urllib necessari requir',)
```

```
-----
('compil linq queri built sql function queri execut c take way much tim
e use sqlcommand command new sqlcommand queri connect sqldataread reade
r command executeread read data take minut execut c much quicker ssms s
ee googl search someth compiledqueri confus whether still use built sql
function year month day getdat anyon show exampl creat call compil quer
i use built function grate thank advanc',)
```

```
In [35]: #Taking 50K entries to a dataframe.
write_db = 'Processed.db'
if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        preprocessed_data = pd.read_sql_query("""SELECT question, Tags
FROM QuestionsProcessed WHERE TAGS IS NOT NULL""", conn_r)
    conn_r.commit()
    conn_r.close()
```



```
In [36]: preprocessed_data.head()
```

Out[36]:

	question	tags
0	carri exist wp user new buddypress instal made...	buddypress
1	self sign certif work behind apach revers prox...	apache ssl
2	python return max number array error python sc...	python arrays numbers max
3	auto start applic user login mac java program ...	packagemaker autostart
4	xcode ctrl drag edit code snippet xcode ctrl d...	xcode4

```
In [37]: print("number of data points in sample :", preprocessed_data.shape[0])
print("number of dimensions :", preprocessed_data.shape[1])
```

```
number of data points in sample : 49999
number of dimensions : 2
```

## 4. Machine Learning Models

### 4.1 Converting tags for multilabel problems

X	y1	y2	y3	y4
x1	0	1	1	0
x1	1	0	0	0
x1	0	1	0	0

```
In [38]: # binary='true' will give a binary vectorizer
vectorizer = CountVectorizer(tokenizer = lambda x: x.split(), binary='true', max_features=20000, ngram_range=(1,2))
multilabel_y = vectorizer.fit_transform(preprocessed_data['tags'])
```


We will sample the number of tags instead considering all of them (due to limitation of computing power)

```
In [39]: def tags_to_choose(n):
        t = multilabel_y.sum(axis=0).tolist()[0]
        sorted_tags_i = sorted(range(len(t)), key=lambda i: t[i], reverse=True)
        multilabel_yn=multilabel_y[:,sorted_tags_i[:n]]
        return multilabel_yn

        def questions_explained_fn(n):
            multilabel_yn = tags_to_choose(n)
            x= multilabel_yn.sum(axis=1)
            return (np.count_nonzero(x==0))
```

```
In [40]: questions_explained = []
        total_tags=multilabel_y.shape[1]
        total_qs=preprocessed_data.shape[0]
        for i in range(500, total_tags, 100):
            questions_explained.append(np.round(((total_qs-questions_explained_fn(i))/total_qs)*100,3))
```

```
In [41]: fig, ax = plt.subplots()
        ax.plot(questions_explained)
        xlabel = list(500+np.array(range(-50,450,50))*50)
        ax.set_xticklabels(xlabel)
        plt.xlabel("Number of tags")
        plt.ylabel("Number Questions covered partially")
        plt.grid()
        plt.show()
        # you can choose any number of tags based on your computing power, minimum is 50(it covers 90% of the tags)
        print("with ",5500,"tags we are covering ",questions_explained[50],"% of questions")
```



with 5500 tags we are covering 98.496 % of questions

```
In [42]: multilabel_yx = tags_to_choose(5500)
print("number of questions that are not covered :", questions_explained_
_fn(5500),"out of ", total_qs)
```

number of questions that are not covered : 752 out of 49999

```
In [43]: print("Number of tags in sample :", multilabel_y.shape[1])
print("number of tags taken :", multilabel_yx.shape[1],"(",(multilabel_
yx.shape[1]/multilabel_y.shape[1])*100,"%")")
```

Number of tags in sample : 20000  
number of tags taken : 5500 ( 27.500000000000004 %)

**We consider top 15% tags which covers 99% of the questions**

## 4.2 Split the data into test and train (80:20)

```
In [44]: total_size=preprocessed_data.shape[0]
train_size=int(0.80*total_size)

x_train=preprocessed_data.head(train_size)
x_test=preprocessed_data.tail(total_size - train_size)

y_train = multilabel_yx[0:train_size,:]
y_test = multilabel_yx[train_size:total_size,:]
```

```
In [45]: print("Number of data points in train data :", y_train.shape)
print("Number of data points in test data :", y_test.shape)
```

Number of data points in train data : (39999, 5500)  
Number of data points in test data : (10000, 5500)

## 4.3 Featurizing data

```
In [46]: start = datetime.now()
vectorizer = TfidfVectorizer(min_df=0.00009, max_features=200000, smooth
h_idf=True, norm="l2", \
                        tokenizer = lambda x: x.split(), sublinear
_tf=False, ngram_range=(1,2))
x_train_multilabel = vectorizer.fit_transform(x_train['question'])
x_test_multilabel = vectorizer.transform(x_test['question'])
print("Time taken to run this cell :", datetime.now() - start)
```

Time taken to run this cell : 0:00:11.048061

```
In [47]: print("Dimensions of train data X:",x_train_multilabel.shape, "Y :",y_t
rain.shape)
print("Dimensions of test data X:",x_test_multilabel.shape,"Y:",y_test.
shape)
```

Dimensions of train data X: (39999, 94094) Y : (39999, 5500)  
Dimensions of test data X: (10000, 94094) Y: (10000, 5500)

```
In [48]: # https://www.analyticsvidhya.com/blog/2017/08/introduction-to-multi-la
```

```

bel-classification/
#https://stats.stackexchange.com/questions/117796/scikit-multi-label-cl
assification
# classifier = LabelPowerset(GaussianNB())
"""
from skmultilearn.adapt import MLkNN
classifier = MLkNN(k=21)

# train
classifier.fit(x_train_multilabel, y_train)

# predict
predictions = classifier.predict(x_test_multilabel)
print(accuracy_score(y_test, predictions))
print(metrics.f1_score(y_test, predictions, average = 'macro'))
print(metrics.f1_score(y_test, predictions, average = 'micro'))
print(metrics.hamming_loss(y_test, predictions))

"""
# we are getting memory error because the multilearn package
# is trying to convert the data into dense matrix
# -----
-----
#MemoryError                                Traceback (most recent call
last)
#<ipython-input-170-f0e7c7f3e0be> in <module>()
#----> classifier.fit(x_train_multilabel, y_train)

```

Out[48]: "\nfrom skmultilearn.adapt import MLkNN\nclassifier = MLkNN(k=21)\n\n# train\nclassifier.fit(x\_train\_multilabel, y\_train)\n\n# predict\npredictions = classifier.predict(x\_test\_multilabel)\nprint(accuracy\_score(y\_test, predictions))\nprint(metrics.f1\_score(y\_test, predictions, average = 'macro'))\nprint(metrics.f1\_score(y\_test, predictions, average = 'micro'))\nprint(metrics.hamming\_loss(y\_test, predictions))\n\n"

## 4.4 Applying Logistic Regression with OneVsRest Classifier

```
In [49]: # this will be taking so much time try not to run it, download the lr_w
ith_equal_weight.pkl file and use to predict
# This takes about 6-7 hours to run.
start = datetime.now()
classifier = OneVsRestClassifier(SGDClassifier(loss='log', alpha=0.0000
1, penalty='l1'), n_jobs=1)
classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict(x_test_multilabel)

print("Accuracy :",metrics.accuracy_score(y_test,predictions))
print("Macro F1 score :",metrics.f1_score(y_test, predictions, average
= 'macro'))
print("Micro F1 score :",metrics.f1_score(y_test, predictions, average
= 'micro'))
print("Hamming Loss :",metrics.hamming_loss(y_test,predictions))

report = metrics.classification_report(y_test, predictions,output_dict=
True)
ovr_classification = pd.DataFrame(report).transpose()
print(" ")

#Printing only the Top 10 Rows out of 5500 in Total
print("Precision recall report :\n")
print(ovr_classification.head(10))
print(" ")

print("Time taken to run this cell :", datetime.now() - start)
```

```
Accuracy : 0.0713
Macro F1 score : 0.06971582995764285
Micro F1 score : 0.32809117080250294
Hamming Loss : 0.0005134727272727272
```

```
Precision recall report :
```

	f1-score	precision	recall	support
0	0.324421	0.549849	0.230088	791.0
1	0.552147	0.773956	0.429155	734.0
10	0.709163	0.847619	0.609589	292.0
100	0.468750	0.681818	0.357143	42.0

1000	0.000000	0.000000	0.000000	2.0
1001	0.000000	0.000000	0.000000	3.0
1002	0.000000	0.000000	0.000000	2.0
1003	0.000000	0.000000	0.000000	4.0
1004	0.000000	0.000000	0.000000	10.0
1005	0.000000	0.000000	0.000000	5.0

Time taken to run this cell : 0:28:27.868461

```
In [50]: #joblib_in = open("lr_with_equal_weight.pkl","rb")
#classifier = joblib.load(joblib_in)
```

## 4.5 Modeling with less data points (50K data points) and more weight to title and 500 tags only.

```
In [51]: sql_create_table = """CREATE TABLE IF NOT EXISTS QuestionsProcessed (qu
estion text NOT NULL, code text, tags text, words_pre integer, words_po
st integer, is_code integer);"""
create_database_table("Titlmoreweight.db", sql_create_table)
```

Tables in the databse:  
QuestionsProcessed

```
In [52]: # http://www.sqlitetutorial.net/sqlite-delete/
# https://stackoverflow.com/questions/2279706/select-random-row-from-a-
sqlite-table

read_db = 'train_no_dup.db'
write_db = 'Titlmoreweight.db'
train_datasize = 40000
if os.path.isfile(read_db):
    conn_r = create_connection(read_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        # for selecting first 50K rows
        reader.execute("SELECT Title, Body, Tags From no_dup_train LIMIT
50001;")
```

```

        # for selecting random points
        #reader.execute("SELECT Title, Body, Tags From no_dup_train ORDER BY RANDOM() LIMIT 500001;")

if os.path.isfile(write_db):
    conn_w = create_connection(write_db)
    if conn_w is not None:
        tables = checkTableExists(conn_w)
        writer = conn_w.cursor()
        if tables != 0:
            writer.execute("DELETE FROM QuestionsProcessed WHERE 1")
            print("Cleared All the rows")

```

Tables in the database:  
 QuestionsProcessed  
 Cleared All the rows

### 4.5.1 Preprocessing of questions

1. Separate Code from Body
2. Remove Special characters from Question title and description (not in code)
3. **Give more weightage to title : Add title three times to the question**
4. Remove stop words (Except 'C')
5. Remove HTML Tags
6. Convert all the characters into small letters
7. Use SnowballStemmer to stem the words

In [53]: `#http://www.bernzilla.com/2008/05/13/selecting-a-random-row-from-an-sqlite-table/`

```

start = datetime.now()
preprocessed_data_list=[]
reader.fetchone()
questions_with_code=0
len_pre=0
len_post=0

```



```

questions_proccesed = 0
for row in reader:

    is_code = 0

    title, question, tags = row[0], row[1], str(row[2])

    if '<code>' in question:
        questions_with_code+=1
        is_code = 1
    x = len(question)+len(title)
    len_pre+=x

    code = str(re.findall(r'<code>(.*?)</code>', question, flags=re.DOTALL))

    question=re.sub('<code>(.*?)</code>', '', question, flags=re.MULTILINE|re.DOTALL)
    question=stripthtml(question.encode('utf-8'))

    title=title.encode('utf-8')

    # adding title three time to the data to increase its weight
    # add tags string to the training data

    question=str(title)+" "+str(title)+" "+str(title)+" "+question

    #     if questions_proccesed<=train_datasize:
    #         question=str(title)+" "+str(title)+" "+str(title)+" "+question
    #     n+" "+str(tags)
    #     else:
    #         question=str(title)+" "+str(title)+" "+str(title)+" "+question
    #     n

    question=re.sub(r'^A-Za-z0-9#+.\-]+', '', question)
    words=word_tokenize(str(question.lower()))

    #Removing all single letter and and stopwords from question exceptt
    for the letter 'c'

```

```

        question=' '.join(str(stemmer.stem(j)) for j in words if j not in s
top_words and (len(j)!=1 or j=='c'))

        len_post+=len(question)
        tup = (question,code,tags,x,len(question),is_code)
        questions_proccesed += 1
        writer.execute("insert into QuestionsProcessed(question,code,tags,w
ords_pre,words_post,is_code) values (?, ?, ?, ?, ?, ?)", tup)
        if (questions_proccesed%100000==0):
            print("number of questions completed=",questions_proccesed)

no_dup_avg_len_pre=(len_pre*1.0)/questions_proccesed
no_dup_avg_len_post=(len_post*1.0)/questions_proccesed

print( "Avg. length of questions(Title+Body) before processing: %d"%no_
dup_avg_len_pre)
print( "Avg. length of questions(Title+Body) after processing: %d"%no_d
up_avg_len_post)
print( "Percent of questions containing code: %d"%((questions_with_code
*100.0)/questions_proccesed))

print("Time taken to run this cell :", datetime.now() - start)

```

```

Avg. length of questions(Title+Body) before processing: 1224
Avg. length of questions(Title+Body) after processing: 437
Percent of questions containing code: 56
Time taken to run this cell : 0:01:56.713871

```

```

In [54]: # never forget to close the conections or else we will end up with data
base locks
conn_r.commit()
conn_w.commit()
conn_r.close()
conn_w.close()

```

#### Sample quesitons after preprocessing of data

```

In [55]: if os.path.isfile(write_db):

```

```

conn_r = create_connection(write_db)
if conn_r is not None:
    reader = conn_r.cursor()
    reader.execute("SELECT question From QuestionsProcessed LIMIT 1
0")
    print("Questions after preprocessed")
    print('='*100)
    reader.fetchone()
    for row in reader:
        print(row)
        print('-'*100)
conn_r.commit()
conn_r.close()

```

Questions after preprocessed

```

=====
=====
('dynam datagrid bind silverlight dynam datagrid bind silverlight dynam
datagrid bind silverlight bind datagrid dynam code wrote code debug cod
e block seem bind correct grid come column form come grid column althou
gh necessari bind nthank repli advance..',)
-----
-----
('java.lang.noclassdeffounderror javax servlet jsp tagext taglibraryval
id java.lang.noclassdeffounderror javax servlet jsp tagext taglibraryva
lid java.lang.noclassdeffounderror javax servlet jsp tagext taglibraryv
alid follow guid link instal jstl got follow error tri launch jsp page
java.lang.noclassdeffounderror javax servlet jsp tagext taglibraryvalid
taglib declar instal jstl 1.1 tomcat webapp tri project work also tri v
ersion 1.2 jstl still messag caus solv',)
-----
-----
('java.sql.sqllexcept microsoft odbc driver manag invalid descriptor ind
ex java.sql.sqllexcept microsoft odbc driver manag invalid descriptor in
dex java.sql.sqllexcept microsoft odbc driver manag invalid descriptor i
ndex use follow code display caus solv',)
-----
-----
('better way updat feed fb php sdk better way updat feed fb php sdk bet
ter way updat feed fb php sdk novic facebook api read mani tutori still

```

```

confused.i find post feed api method like correct second way use curl s
ometh like way better',)
-----
('btnadd click event open two window record ad btnadd click event open
two window record ad btnadd click event open two window record ad open
window search.aspx use code hav add button search.aspx nwhen insert rec
ord btnadd click event open anoth window nafter insert record close win
dow',)
-----
('sql inject issu prevent correct form submit php sql inject issu prev
ent correct form submit php sql inject issu prevent correct form submi
ss php check everyth think make sure input field safe type sql inject g
ood news safe bad news one tag mess form submit place even touch life
figur exact html use templat file forgiv okay entir php script get exec
ut see data post none forum field post problem use someth titl field no
ne data get post current use print post see submit noth work flawless s
tatement though also mention script work flawless local machin use host
come across problem state list input test mess',)
-----
('countabl subaddit lebesgu measur countabl subaddit lebesgu measur cou
ntabl subaddit lebesgu measur let lbrace rbrace sequenc set sigma -alge
bra mathcal want show left bigcup right leq sum left right countabl add
it measur defin set sigma algebra mathcal think use monoton properti so
mewher proof start appreci littl help nthank ad han answer make follow
addit construct given han answer clear bigcup bigcup cap emptyset neq l
eft bigcup right left bigcup right sum left right also construct subset
monoton left right leq left right final would sum leq sum result follo
w',)
-----
('hql equival sql queri hql equival sql queri hql equival sql queri hql
queri replac name class properti name error occur hql error',)
-----
('undefin symbol architectur i386 objc class skpsmtpmessag referenc err
or undefin symbol architectur i386 objc class skpsmtpmessag referenc er
ror undefin symbol architectur i386 objc class skpsmtpmessag referenc e

```

```

for underlin symbol architectur 1500 obje class skpsmtpmessag referenc e
rror import framework send email applic background import framework i.e
skpsmtpmessag somebody suggest get error collect2 ld return exit status
import framework correct sorc taken framework follow mfmcomposeviewc
ontrol question lock field updat answer drag drop folder project click
copi nthat',)
-----
-----

```

### Saving Preprocessed data to a Database

```

In [56]: #Taking 50K entries to a dataframe.
write_db = 'Titlmoreweight.db'
if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        preprocessed_data = pd.read_sql_query("""SELECT question, Tags
FROM QuestionsProcessed""", conn_r)
    conn_r.commit()
    conn_r.close()

```

```

In [57]: preprocessed_data.head()

```

Out[57]:

	question	tags
0	dynam datagrid bind silverlight dynam datagrid...	c# silverlight data-binding
1	dynam datagrid bind silverlight dynam datagrid...	c# silverlight data-binding columns
2	java.lang.noclassdeffoundererror javax servlet j...	jsp jstl
3	java.sql.sqlexcept microsoft odbc driver manag...	java jdbc
4	better way updat feed fb php sdk better way up...	facebook api facebook-php-sdk

```

In [58]: print("number of data points in sample :", preprocessed_data.shape[0])
print("number of dimensions :", preprocessed_data.shape[1])

```

```
number of data points in sample : 50000
number of dimensions : 2
```

### Converting string Tags to multilable output variables


```
In [59]: vectorizer = CountVectorizer(tokenizer = lambda x: x.split(), binary='true',max_features=20000,ngram_range=(1,2))
multilabel_y = vectorizer.fit_transform(preprocessed_data['tags'])
```

### Selecting 500 Tags

```
In [60]: questions_explained = []
total_tags=multilabel_y.shape[1]
total_qs=preprocessed_data.shape[0]
for i in range(500, total_tags, 100):
    questions_explained.append(np.round(((total_qs-questions_explained_
fn(i))/total_qs)*100,3))
```

```
In [61]: fig, ax = plt.subplots()
ax.plot(questions_explained)
xlabel = list(500+np.array(range(-50,450,50))*50)
ax.set_xticklabels(xlabel)
plt.xlabel("Number of tags")
plt.ylabel("Number Questions covered partially")
plt.grid()
plt.show()

# you can choose any number of tags based on your computing power, minimum is 500(it covers 90% of the tags)
print("with ",5500,"tags we are covering ",questions_explained[50],"% of questions")
print("with ",500,"tags we are covering ",questions_explained[0],"% of questions")
```



with 5500 tags we are covering 98.88 % of questions  
 with 500 tags we are covering 90.818 % of questions

```
In [62]: # we will be taking 500 tags
          multilabel_yx = tags_to_choose(500)
          print("number of questions that are not covered :", questions_explained_
                _fn(500),"out of ", total_qs)
```

number of questions that are not covered : 4591 out of 50000

```
In [63]: x_train=preprocessed_data.head(train_datasize)
          x_test=preprocessed_data.tail(preprocessed_data.shape[0] - 40000)

          y_train = multilabel_yx[0:train_datasize,:]
          y_test = multilabel_yx[train_datasize:preprocessed_data.shape[0],:]
```

```
In [64]: print("Number of data points in train data :", y_train.shape)
          print("Number of data points in test data :", y_test.shape)
```

Number of data points in train data : (40000, 500)

Number of data points in test data : (10000, 500)

```
In [65]: print(preprocessed_data.shape[0])
```

50000

## Application of TFIDF :-

### 4.5.2 Featurizing Data with TFIDF Vectorizer

```
In [66]: start = datetime.now()
vectorizer = TfidfVectorizer(min_df=0.00009, max_features=20000, smooth_idf=True, norm="l2", \
                             tokenizer = lambda x: x.split(), sublinear_tf=False, ngram_range=(1,2))
x_train_multilabel = vectorizer.fit_transform(x_train['question'])
x_test_multilabel = vectorizer.transform(x_test['question'])
print("Time taken to run this cell :", datetime.now() - start)
```

Time taken to run this cell : 0:00:12.606605

```
In [67]: print("Dimensions of train data X:", x_train_multilabel.shape, "Y :", y_train.shape)
print("Dimensions of test data X:", x_test_multilabel.shape, "Y:", y_test.shape)
```

Dimensions of train data X: (40000, 20000) Y : (40000, 500)

Dimensions of test data X: (10000, 20000) Y: (10000, 500)

### 4.5.3 (TFIDF) Applying Logistic Regression :- OneVsRest Classifier + SGDClassifier (loss='log')

```
In [68]: start = datetime.now()
```



```

classifier = OneVsRestClassifier(SGDClassifier(loss='log', alpha=0.0000
1, penalty='l1'), n_jobs=-1)
classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict(x_test_multilabel)

print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming Loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))
print(" ")

report = metrics.classification_report(y_test, predictions,output_dict=True)
classification = pd.DataFrame(report).transpose()
print(" ")

#Printing only the Top 10 Rows out of 500 in Total
print("Precision recall report :\n")
print(classification.head(10))
print(" ")

print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.207
Hamming Loss  0.0033218
Micro-average quality numbers

```

Precision: 0.6645, Recall: 0.2679, F1-measure: 0.3819

Macro-average quality numbers

Precision: 0.4635, Recall: 0.2007, F1-measure: 0.2610

Precision recall report :

	f1-score	precision	recall	support
0	0.318937	0.623377	0.214286	224.0
1	0.305419	0.642487	0.200323	619.0
10	0.366947	0.623810	0.259921	504.0
100	0.521739	0.750000	0.400000	30.0
101	0.378788	0.714286	0.257732	97.0
102	0.192308	0.312500	0.138889	36.0
103	0.196078	0.333333	0.138889	36.0
104	0.312500	0.500000	0.227273	22.0
105	0.000000	0.000000	0.000000	8.0
106	0.045455	0.111111	0.028571	35.0

Time taken to run this cell : 0:00:24.330834

```
In [69]: #joblib.dump(classifier, 'lr_with_more_title_weight.pkl')
```

#### 4.5.4 (TFIDF) Applying Logistic Regression :- OneVsRest Classifier + Logistic Regression

```
In [70]: start = datetime.now()
classifier_2 = OneVsRestClassifier(LogisticRegression(penalty='l2'), n_
jobs=-1)
classifier_2.fit(x_train_multilabel, y_train)
predictions_2 = classifier_2.predict(x_test_multilabel)
print("Accuracy :",metrics.accuracy_score(y_test, predictions_2))
print("Hamming Loss ",metrics.hamming_loss(y_test,predictions_2))

precision = precision_score(y_test, predictions_2, average='micro')
recall = recall_score(y_test, predictions_2, average='micro')
f1 = f1_score(y_test, predictions_2, average='micro')
```

```

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

precision = precision_score(y_test, predictions_2, average='macro')
recall = recall_score(y_test, predictions_2, average='macro')
f1 = f1_score(y_test, predictions_2, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

report = metrics.classification_report(y_test, predictions_2, output_dict=True)
classification = pd.DataFrame(report).transpose()
print(" ")

#Printing only the Top 10 Rows out of 500 in Total
print("Precision recall report :\n")
print(classification.head(10))
print(" ")

print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.1596
Hamming Loss 0.0034734
Micro-average quality numbers
Precision: 0.7952, Recall: 0.1253, F1-measure: 0.2165
Macro-average quality numbers
Precision: 0.2431, Recall: 0.0495, F1-measure: 0.0762

```

Precision recall report :

	f1-score	precision	recall	support
0	0.325260	0.723077	0.209821	224.0
1	0.220708	0.704348	0.130856	619.0
10	0.257329	0.718182	0.156746	504.0
100	0.235294	1.000000	0.133333	30.0
101	0.150943	0.888889	0.082474	97.0

102	0.097561	0.400000	0.055556	36.0
103	0.142857	0.500000	0.083333	36.0
104	0.160000	0.666667	0.090909	22.0
105	0.000000	0.000000	0.000000	8.0
106	0.054054	0.500000	0.028571	35.0

Time taken to run this cell : 0:00:41.021962

#### 4.5.5 (TFIDF) Hyperparameter Tuning on "Logistic Regression :- OneVsRest Classifier + SGDClassifier(loss='log')"

```
In [87]: x_cv_multilabel = x_train_multilabel[32000:40000]

start_time = datetime.now()
warnings.filterwarnings('ignore')
params = {"estimator__alpha":C_hyperparam}
#Carrying out 5-fold Cross Validation.

logistic = OneVsRestClassifier(SGDClassifier(loss='log',penalty='l1'),
n_jobs=-1)
TFIDF_model1 = GridSearchCV(logistic,params,scoring='f1_micro', cv=5,n_
jobs=1,verbose=1)
TFIDF_model1.fit(x_train_multilabel,y_train)

print("Time Required to Carry out Hyperparameter Tuning :- ", datetime.
now()-start_time)
print("Best Values Obtained :-", TFIDF_model1.best_estimator_)
print("Best Cross Validation Score :-", TFIDF_model1.best_score_)
```

Fitting 5 folds for each of 7 candidates, totalling 35 fits

[Parallel(n\_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[Parallel(n\_jobs=1)]: Done 35 out of 35 | elapsed: 12.3min finished

Time Required to Carry out Hyperparameter Tuning :- 0:12:44.195745  
Best Values Obtained :- OneVsRestClassifier(estimator=SGDClassifier(alpha=1e-06, average=False,

class\_weight=None

```

on=0.1,
e,
ss='log',
ange=5,
one,
rbose=0,
n_jobs=-1)
Best Cross Validation Score :- 0.40200532332152966

```

```

class_weight=None,
early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
l1_ratio=0.15, learning_rate='optimal', loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None, penalty='l1', power_t=0.5, random_state=None,
shuffle=True, tol=0.001, validation_fraction=0.1, verbose=0, warm_start=False),

```

#### 4.5.6 Applying Optimized Logistic Regression :- "SGDClassifier(loss='log') + OneVsRest Classifier" on Test Dataset

```

In [88]: classifier = OneVsRestClassifier(estimator=SGDClassifier(alpha=10**(-6), average=False, class_weight=None,
early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True, l1_ratio=0.15, learning_rate='optimal',
loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None, penalty='l1', power_t=0.5, random_state=None,
shuffle=True, tol=0.001, validation_fraction=0.1, verbose=0, warm_start=False), n_jobs=-1)

```

```

In [89]: start = datetime.now()

classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict(x_test_multilabel)

```

```

print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test,predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1= f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision,recall,f1))

precision = precision_score(y_test,predictions, average='macro')
recall = recall_score(y_test,predictions, average='macro')
f1 = f1_score(y_test,predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

report = metrics.classification_report(y_test, predictions,output_dict=True)
classification = pd.DataFrame(report).transpose()
print(" ")

#Printing only the Top 10 Rows out of 500 in Total
print("Precision recall report :\n")
print(classification.head(10))
print(" ")

print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.1815
Hamming loss  0.0036454
Micro-average quality numbers
Precision: 0.5426, Recall: 0.3070, F1-measure: 0.3921
Macro-average quality numbers
Precision: 0.4277, Recall: 0.2226, F1-measure: 0.2742

Precision recall report :

```

	f1-score	precision	recall	support
0	0.278481	0.264000	0.294643	224.0
1	0.357333	0.397233	0.324717	619.0
10	0.368742	0.479365	0.299603	504.0
100	0.521739	0.750000	0.400000	30.0
101	0.342857	0.558140	0.247423	97.0
102	0.222222	0.333333	0.166667	36.0
103	0.192308	0.312500	0.138889	36.0
104	0.352941	0.500000	0.272727	22.0
105	0.000000	0.000000	0.000000	8.0
106	0.090909	0.222222	0.057143	35.0

Time taken to run this cell : 0:00:29.102101

## Application of BOW - Upto Bi-Grams :-

### 5.1 Featurizing data with BOW Vectorizer

```
In [73]: start = datetime.now()
BOW_vectorizer = CountVectorizer(min_df=0.00009,max_features=20000,
                                tokenizer = lambda x: x.split(), ngram
                                _range=(1,2))
x_bow_train_multilabel = BOW_vectorizer.fit_transform(x_train['question'])
x_bow_test_multilabel = BOW_vectorizer.transform(x_test['question'])
print("Time taken to run this cell :", datetime.now() - start)
```

Time taken to run this cell : 0:00:13.971770

```
In [74]: print("Dimensions of train data X:",x_bow_train_multilabel.shape, "Y : "
,y_train.shape)
print("Dimensions of test data X:",x_bow_test_multilabel.shape,"Y:",y_test.shape)
```

Dimensions of train data X: (40000, 20000) Y : (40000, 500)  
Dimensions of test data X: (10000, 20000) Y: (10000, 500)

## 5.2 (BOW) Applying Logistic Regression :- OneVsRest Classifier + SGDClassifier (loss='log')

```
In [75]: start = datetime.now()
classifier2 = OneVsRestClassifier(SGDClassifier(loss='log', alpha=0.00001, penalty='l1'), n_jobs=1)
classifier2.fit(x_bow_train_multilabel, y_train)
predictions2 = classifier2.predict(x_bow_test_multilabel)

print("Accuracy :", metrics.accuracy_score(y_test, predictions2))
print("Hamming loss ", metrics.hamming_loss(y_test, predictions2))

precision2 = precision_score(y_test, predictions2, average='micro')
recall2 = recall_score(y_test, predictions2, average='micro')
f1_2 = f1_score(y_test, predictions2, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision2, recall2, f1_2))

precision2 = precision_score(y_test, predictions2, average='macro')
recall2 = recall_score(y_test, predictions2, average='macro')
f1_2 = f1_score(y_test, predictions2, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision2, recall2, f1_2))

report = metrics.classification_report(y_test, predictions2, output_dict=True)
classification = pd.DataFrame(report).transpose()
print(" ")

#Printing only the Top 10 Rows out of 500 in Total
print("Precision recall report :\n")
print(classification.head(10))
```



```
print(" ")
print("Time taken to run this cell :", datetime.now() - start)
```

```
Accuracy : 0.0645
Hamming loss 0.0098328
Micro-average quality numbers
Precision: 0.1601, Recall: 0.3690, F1-measure: 0.2233
Macro-average quality numbers
Precision: 0.1102, Recall: 0.2771, F1-measure: 0.1435
```

Precision recall report :

	f1-score	precision	recall	support
0	0.198758	0.137694	0.357143	224.0
1	0.339324	0.305699	0.381260	619.0
10	0.347261	0.326353	0.371032	504.0
100	0.158273	0.100917	0.366667	30.0
101	0.372093	0.426667	0.329897	97.0
102	0.115226	0.067633	0.388889	36.0
103	0.136646	0.088000	0.305556	36.0
104	0.155340	0.098765	0.363636	22.0
105	0.000000	0.000000	0.000000	8.0
106	0.065574	0.040541	0.171429	35.0

Time taken to run this cell : 0:17:20.428468

### 5.3 (BOW) Applying Logistic Regression :- OneVsRest Classifier + Logistic Regression

```
In [76]: start = datetime.now()
classifier3 = OneVsRestClassifier(LogisticRegression(C=1.0,penalty='l2'
), n_jobs=1)
classifier3.fit(x_bow_train_multilabel, y_train)
predictions3 = classifier3.predict (x_bow_test_multilabel)

print("Accuracy :",metrics.accuracy_score(y_test, predictions3))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions3))
```

```

precision3 = precision_score(y_test, predictions3, average='micro')
recall3 = recall_score(y_test, predictions3, average='micro')
f1_3= f1_score(y_test, predictions3, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision3, recall3, f1_3))

precision3 = precision_score(y_test, predictions3, average='macro')
recall3 = recall_score(y_test, predictions3, average='macro')
f1_3 = f1_score(y_test, predictions3, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision3, recall3, f1_3))

report = metrics.classification_report(y_test, predictions3, output_dict=True)
classification = pd.DataFrame(report).transpose()
print(" ")

#Printing only the Top 10 Rows out of 500 in Total
print("Precision recall report :\n")
print(classification.head(10))
print(" ")

print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.1707
Hamming loss  0.0037462
Micro-average quality numbers
Precision: 0.5206, Recall: 0.2764, F1-measure: 0.3611
Macro-average quality numbers
Precision: 0.3819, Recall: 0.1916, F1-measure: 0.2401

```

Precision recall report :

	f1-score	precision	recall	support
0	0.289044	0.302439	0.276786	224.0

1	0.367424	0.443936	0.313409	619.0
10	0.344473	0.489051	0.265873	504.0
100	0.474576	0.482759	0.466667	30.0
101	0.302158	0.500000	0.216495	97.0
102	0.233333	0.291667	0.194444	36.0
103	0.218182	0.315789	0.166667	36.0
104	0.193548	0.333333	0.136364	22.0
105	0.000000	0.000000	0.000000	8.0
106	0.000000	0.000000	0.000000	35.0

Time taken to run this cell : 0:29:07.881966

## 5.4 (BOW) Hyperparameter Tuning on "Logistic Regression :- OneVsRest Classifier + Logistic Regression"

```
In [77]: C_hyperparam =[]
#initializing an empty list

for a in range(-6,7,2):
    C_hyperparam.append(10**a)

print(C_hyperparam)
```

```
[1e-06, 0.0001, 0.01, 1, 100, 10000, 1000000]
```

```
In [78]: from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import StandardScaler

#performing scaling
Scaler1 = StandardScaler(with_mean=False)
X_Train_SCBOW = Scaler1.fit_transform(x_bow_train_multilabel)
X_Test_SCBOW = Scaler1.transform(x_bow_test_multilabel)
```

Basically to carry out the Hyperparameter Tuning with the CV Dataset, we first need to obtain the CV Data :- We will be carrying out Simple CV instead of GridSearchCV so as to minimise the Time Complexity for the same.

We will obtain the CV Dataset by splitting the Train Dataset :- x\_bow\_train\_multilabel in the 80:20 ratio.

```
In [79]: #Obtaining the CV Dataset from the Train Dataset :- X_CV
x_bow_cv_multilabel = x_bow_train_multilabel[32000:40000]
X_CV_SCBOW = Scaler1.transform(x_bow_cv_multilabel)

print(X_CV_SCBOW.shape)

(8000, 20000)
```

```
In [80]: #Obtaining the CV Dataset from the Train Dataset :- Y_CV
y_cv = y_train[32000:40000]
print(y_cv.shape)

(8000, 500)
```

### **(BOW) Hyperparameter Tuning using GridSearch on Better Model :-**

```
In [82]: start_time = datetime.now()
warnings.filterwarnings('ignore')
params = {"estimator__C":C_hyperparam}

#Carrying out 5-fold Cross Validation.

logistic = OneVsRestClassifier(LogisticRegression(penalty='l2'), n_jobs=1)
BOW_model1 = GridSearchCV(logistic,params,scoring='f1_micro', cv=5,n_jobs=1,verbose=1)
BOW_model1.fit(X_Train_SCBOW,y_train)

print("Time Required to Carry out Hyperparameter Tuning :- ", datetime.now()-start_time)
print("Best Values Obtained :-", BOW_model1.best_estimator_)
print("Best Cross Validation Score :-", BOW_model1.best_score_)
```

Fitting 5 folds for each of 7 candidates, totalling 35 fits

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
```

```
[Parallel(n_jobs=1)]: Done 35 out of 35 | elapsed: 683.6min finished
```

```
Time Required to Carry out Hyperparameter Tuning :- 12:02:43.933640
Best Values Obtained :- OneVsRestClassifier(estimator=LogisticRegression(C=100, class_weight=None, dual=False, fit_intercept=True, intercept_scaling=1, l1_ratio=None, max_iter=100, multi_class='auto', n_jobs=None, penalty='l2', random_state=None, solver='lbfgs', tol=0.0001, verbose=0, warm_start=False), n_jobs=1)
Best Cross Validation Score :- 0.307864297932658
```

## 5.5 Applying Optimized Logistic Regression :- "Logistic Regression + OneVsRest Classifier" on Test Dataset

```
In [81]: BOW_classifier = OneVsRestClassifier(LogisticRegression(C=100, class_weight=None, dual=False, fit_intercept=True, intercept_scaling=1, l1_ratio=None, max_iter=100, multi_class='auto', n_jobs=None, penalty='l2', random_state=None, solver='lbfgs', tol=0.0001, verbose=0, warm_start=False), n_jobs=1)
```

```
In [82]: start = datetime.now()
BOW_classifier = OneVsRestClassifier(LogisticRegression(C=100, penalty='l2'), n_jobs=1)
```

```

BOW_classifier.fit(x_bow_train_multilabel, y_train)
BOW_predictions = BOW_classifier.predict(x_bow_test_multilabel)

print("Accuracy :",metrics.accuracy_score(y_test, BOW_predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,BOW_predictions))

BOW_precision = precision_score(y_test, BOW_predictions, average='micro')
BOW_recall = recall_score(y_test, BOW_predictions, average='micro')
BOW_f1= f1_score(y_test, BOW_predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(BOW_precision, BOW_recall, BOW_f1))

BOW_precision = precision_score(y_test,BOW_predictions, average='macro')
BOW_recall = recall_score(y_test,BOW_predictions, average='macro')
BOW_f1 = f1_score(y_test,BOW_predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(BOW_precision, BOW_recall, BOW_f1))

report = metrics.classification_report(y_test,BOW_predictions,output_dict=True)
classification = pd.DataFrame(report).transpose()
print(" ")

#Printing only the Top 10 Rows out of 500 in Total
print("Precision recall report :\n")
print(classification.head(10))
print(" ")

print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.1498
Hamming loss 0.004091
Micro-average quality numbers
Precision: 0.4489, Recall: 0.2994, F1-measure: 0.3592

```

Macro-average quality numbers  
Precision: 0.3543, Recall: 0.2098, F1-measure: 0.2482

Precision recall report :

	f1-score	precision	recall	support
0	0.240708	0.199413	0.303571	224.0
1	0.353511	0.353226	0.353796	619.0
10	0.357820	0.444118	0.299603	504.0
100	0.451613	0.437500	0.466667	30.0
101	0.315068	0.469388	0.237113	97.0
102	0.253968	0.296296	0.222222	36.0
103	0.245614	0.333333	0.194444	36.0
104	0.294118	0.416667	0.227273	22.0
105	0.000000	0.000000	0.000000	8.0
106	0.000000	0.000000	0.000000	35.0

Time taken to run this cell : 0:33:49.036671

## Application of Linear SVM :-

### 5.6 (BOW) Applying Linear SVM :- OneVsRest Classifier + SGDClassifier (loss='hinge')

```
In [83]: start = datetime.now()
classifier4 = OneVsRestClassifier(SGDClassifier(loss='hinge', alpha=0.0001, penalty='l1'), n_jobs=1)
classifier4.fit(x_train_multilabel, y_train)
predictions4 = classifier4.predict(x_test_multilabel)

print("Accuracy :", metrics.accuracy_score(y_test, predictions4))
print("Hamming loss ", metrics.hamming_loss(y_test, predictions4))

precision4 = precision_score(y_test, predictions4, average='micro')
recall4 = recall_score(y_test, predictions4, average='micro')
f1_4 = f1_score(y_test, predictions4, average='micro')
```

```

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision4, recall4, f1_4))

precision4 = precision_score(y_test, predictions4, average='macro')
recall4 = recall_score(y_test, predictions4, average='macro')
f1_4 = f1_score(y_test, predictions4, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision4, recall4, f1_4))

report = metrics.classification_report(y_test, predictions4, output_dict=True)
classification = pd.DataFrame(report).transpose()
print(" ")

#Printing only the Top 10 Rows out of 500 in Total
print("Precision recall report :\n")
print(classification.head(10))
print(" ")

print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.2122
Hamming loss 0.0032816
Micro-average quality numbers
Precision: 0.6765, Recall: 0.2744, F1-measure: 0.3904
Macro-average quality numbers
Precision: 0.3895, Recall: 0.1989, F1-measure: 0.2441

```

Precision recall report :

	f1-score	precision	recall	support
0	0.321637	0.466102	0.245536	224.0
1	0.345652	0.528239	0.256866	619.0
10	0.323319	0.579487	0.224206	504.0
100	0.500000	0.666667	0.400000	30.0
101	0.358209	0.648649	0.247423	97.0
---	-----	-----	-----	---



```
102 0.000000 0.000000 0.000000 36.0
103 0.054054 1.000000 0.027778 36.0
104 0.275862 0.571429 0.181818 22.0
105 0.000000 0.000000 0.000000 8.0
106 0.000000 0.000000 0.000000 35.0
```

Time taken to run this cell : 0:01:56.350026

## 5.7 (BOW) Applying Linear SVM :- OneVsRest Classifier + Linear SVC

```
In [84]: from sklearn.svm import LinearSVC

start = datetime.now()
classifier5 = OneVsRestClassifier(LinearSVC(loss='squared_hinge', C=1.0,
    penalty='l2'), n_jobs=1)
classifier5.fit(x_train_multilabel, y_train)
predictions5 = classifier5.predict(x_test_multilabel)

print("Accuracy :", metrics.accuracy_score(y_test, predictions5))
print("Hamming loss ", metrics.hamming_loss(y_test, predictions5))

precision5 = precision_score(y_test, predictions5, average='micro')
recall5 = recall_score(y_test, predictions5, average='micro')
f1_5 = f1_score(y_test, predictions5, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision5, recall5, f1_5))

precision5 = precision_score(y_test, predictions5, average='macro')
recall5 = recall_score(y_test, predictions5, average='macro')
f1_5 = f1_score(y_test, predictions5, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision5, recall5, f1_5))

report = metrics.classification_report(y_test, predictions5, output_dict
```

```

=True)
classification = pd.DataFrame(report).transpose()
print(" ")

#Printing only the Top 10 Rows out of 500 in Total
print("Precision recall report :\n")
print(classification.head(10))
print(" ")

print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.2114
Hamming loss 0.0032714
Micro-average quality numbers
Precision: 0.7029, Recall: 0.2525, F1-measure: 0.3716
Macro-average quality numbers
Precision: 0.4377, Recall: 0.1723, F1-measure: 0.2292

```

Precision recall report :

	f1-score	precision	recall	support
0	0.364198	0.590000	0.263393	224.0
1	0.347144	0.565693	0.250404	619.0
10	0.356125	0.631313	0.248016	504.0
100	0.511628	0.846154	0.366667	30.0
101	0.343750	0.709677	0.226804	97.0
102	0.095238	0.333333	0.055556	36.0
103	0.097561	0.400000	0.055556	36.0
104	0.214286	0.500000	0.136364	22.0
105	0.000000	0.000000	0.000000	8.0
106	0.000000	0.000000	0.000000	35.0

Time taken to run this cell : 0:01:17.783593

## 5.8 (BOW) Hyperparameter Tuning on "Linear SVM :- OneVsRest Classifier + Linear SVC"

In [96]: `from sklearn.model_selection import GridSearchCV`

```

start_time = datetime.now()
warnings.filterwarnings('ignore')
params = {"estimator__alpha":C_hyperparam}

#Carrying out 5-fold Cross Validation.

svm = OneVsRestClassifier(SGDClassifier(loss='hinge', penalty='l1'), n_
jobs=1)
svm_model1 = GridSearchCV(svm,params,scoring='f1_micro', cv=5,n_jobs=1,
verbose=1)
svm_model1.fit(x_train_multilabel,y_train)

print("Time Required to Carry out Hyperparameter Tuning :- ", datetime.
now()-start_time)
print("Best Values Obtained :-", svm_model1.best_estimator_)
print("Best Cross Validation Score :-", svm_model1.best_score_)

```

Fitting 5 folds for each of 7 candidates, totalling 35 fits

[Parallel(n\_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[Parallel(n\_jobs=1)]: Done 35 out of 35 | elapsed: 59.0min finished

Time Required to Carry out Hyperparameter Tuning :- 1:01:42.124050  
Best Values Obtained :- OneVsRestClassifier(estimator=SGDClassifier(alp  
ha=1e-06, average=False,

```

                                class_weight=None,
                                early_stopping=False, epsil
on=0.1,
                                eta0=0.0, fit_intercept=True
e,
                                l1_ratio=0.15,
                                learning_rate='optimal',
                                loss='hinge', max_iter=100
0,
                                n_iter_no_change=5, n_jobs=
None,
                                penalty='l1', power_t=0.5,
                                random_state=None, shuffle=
True

```

```

True,
                                tol=0.001, validation_fract
ion=0.1,

                                verbose=0, warm_start=False),
                                n_jobs=1)
Best Cross Validation Score :- 0.3914180794669484

```

## 5.9 Applying Optimized Linear SVM :- "OneVsRest Classifier + SGDClassifier" on Test Dataset

```

In [85]: from sklearn.svm import LinearSVC
SVM_classifier = OneVsRestClassifier(SGDClassifier(alpha=10**(-6), average=False,
                                class_weight=None, early_stopping=False, epsilon=0.1, eta0=0.0,
                                fit_intercept=True, l1_ratio=0.15,
                                learning_rate='optimal', loss='hinge', max_iter=1000, n_iter_no_change=5,
                                n_jobs=None, penalty='l1',
                                power_t=0.5, random_state=None, shuffle=True, tol=0.001,
                                validation_fraction=0.1,
                                verbose=0, warm_start=False))

```

```

In [86]: start = datetime.now()

SVM_classifier.fit(x_train_multilabel, y_train)
SVM_predictions = SVM_classifier.predict(x_test_multilabel)

print("Accuracy :", metrics.accuracy_score(y_test, SVM_predictions))
print("Hamming loss ", metrics.hamming_loss(y_test, SVM_predictions))

SVM_precision = precision_score(y_test, SVM_predictions, average='micro')
SVM_recall = recall_score(y_test, SVM_predictions, average='micro')
SVM_f1 = f1_score(y_test, SVM_predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(SV

```

```

M_precision, SVM_recall, SVM_f1))

SVM_precision = precision_score(y_test,SVM_predictions, average='macro')
SVM_recall = recall_score(y_test,SVM_predictions, average='macro')
SVM_f1 = f1_score(y_test,SVM_predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(SVM_precision, SVM_recall, SVM_f1))

report = metrics.classification_report(y_test, SVM_predictions,output_dict=True)
classification = pd.DataFrame(report).transpose()
print(" ")

#Printing only the Top 10 Rows out of 500 in Total
print("Precision recall report :\n")
print(classification.head(10))
print(" ")

print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.1495
Hamming loss 0.004062
Micro-average quality numbers
Precision: 0.4556, Recall: 0.3115, F1-measure: 0.3700
Macro-average quality numbers
Precision: 0.3453, Recall: 0.2310, F1-measure: 0.2618

```

Precision recall report :

	f1-score	precision	recall	support
0	0.225914	0.179894	0.303571	224.0
1	0.343928	0.347039	0.340872	619.0
10	0.350792	0.454259	0.285714	504.0
100	0.461538	0.545455	0.400000	30.0
101	0.326241	0.522727	0.237113	97.0
102	0.281250	0.321429	0.250000	36.0
103	0.107143	0.150000	0.083333	36.0

104	0.324324	0.400000	0.272727	22.0
105	0.000000	0.000000	0.000000	8.0
106	0.000000	0.000000	0.000000	35.0

Time taken to run this cell : 0:02:55.102056

## 6. Conclusions :-

- StackOverflow is a Website where Users can post Programming Questions for other users to answer the same. The Objective in the case of StackOverflow Tag Prediction is to predict tags for the Questions that are posted on StackOverflow. There are a bunch of different domains on StackOverflow and it is important that the questions are given labels in a way that maximises User Experience. Eg :- If the following question is asked :- "What is the Difference between Pointers and References in C++?". Now StackOverFlow should give Tags such as "C++", "Pointers" etc. This helps the whole community because the people who have followed these Topics and are Programmers themselves in these languages can help the other person out. It also reduces the case that a Question related to Python is sent to a C Programmer with no language on the intricacies of Python.
- Note that a Single Question over here could have Multiple Tags assigned to it. This means that it is not just a Multiclass Classification Problem but a Multilabel Classification Problem. Also following are the constraints that we are dealing with while trying to solve this problem :- Incorrect Tags could Impact the User Experience on StackOverFlow. There are no strict Latency Constraints.

Based on these 2 Conditions we can assume that we Require a Model with High Precision and Recall Values, and the metric that we tried to optimize was the Micro-F1 Score :- weighted average of F1 Score across all the Classes that we had. This is a very good metric, especially when we have Class Imbalance. (We saw this when we carried out EDA in the Beginning).

- We have a total of approximately 6M Questions. Each question over here has the following features - 'Title','Body','Tags'. However, in the interest of time and the Computational Limitations of my Laptop, I am loading only the Top 0.5M rows into a database, and I will

sample only 50K Datapoints. {Only for 50K Datapoints, the Hyperparameter Tuning in the case of Logistic Regression Takes more than 12 Hours}.

- We carry out EDA to obtain information such as which are the most Frequent Tags appearing across the questions and find that "c#,java,php" are some of the most frequent Tags on the Platform and we had seen that there are a total of approx. 20K Unique Tags, out of which 98% variance is retained when we use 5500 Tags and 90% variance is retained when we have 500 Tags. Again, in the interest of Time and Low Computational Power available, we will only be taking 500 Tags for our analysis. After this stage preprocessing of questions is carried out where code snippets are separated and special characters are removed from both the Title as well as the Body. All the stopwords are removed followed by Stemming.
- We featurize our models with the help of BOW(1,2 grams) and TFIDF Vectorizers, and for this we will sample 50K Data Points with 3 times more weight added to the titles. {Because usually Titles have more keywords than the code or Body}. Note that we have used OneVSRestClassifier with both SGDClassifier Logistic Regression Computation as well as with Simple Logistic Regression Computation, because of the algorithmic differences in how these 2 Values are computed.
- Finally we can compare all of our models where we realize that the implementation of "Linear SVM with SGDClassifier (loss='hinge')" using TFIDF Vectorizer resulted in us having the best Micro F1-Score of 0.394 which is very close to our second best model with a Micro-F1 Score 0.393. {Logistic Regression:OVR + SGDClassifier- Tuned using TFIDF Vectorization}

Please take a look at the Different Models that we tried to Train and the metric values that we achieved in relation to Accuracy, Precision, Recall, Hamming Loss, The Best Hyperparameter obtained after carrying out GridSearch for the values of lambda or alpha and the Micro F1-Scores thus obtained.

```
In [112]: from prettytable import PrettyTable

print("Values Obtained after working with 50K Datapoints :-")
print(" ")
```

```

x=PrettyTable()
x.field_names=["Model","Vectorizer","Param","Acc","Hamming","Pre","Re",
"Mic_F1"]

x.add_row(["Log Regression:OVR + SGDClassifier(loss='log')","TFIDF","a=
0.01","0.21","0.003","0.67","0.27","0.38"])
x.add_row(["Log Regression:OVR + LogisticRegression","TFIDF","-","0.16"
,"0.003","0.80","0.16","0.22"])
x.add_row(["Log Regression:OVR + SGDClassifier- Tuned","TFIDF","a=10^(-
6)","0.18","0.003","0.55","0.31","0.393"])
x.add_row(["Log Regression:OVR + SGDClassifier(loss='log')","BOW","a=
0.01","0.06","0.009","0.16","0.36","0.22"])
x.add_row(["Log Regression:OVR + LogisticRegression","BOW","C=1.0","0.
17","0.003","0.52","0.28","0.360"])
x.add_row(["Log Regression:OVR + LogisticRegression - Tuned","BOW","C=
100","0.14","0.004","0.45","0.30","0.361"])
x.add_row(["Linear SVM:OVR + SGDClassifier(loss='hinge')","TFIDF","a=0.
01","0.21","0.003","0.68","0.28","0.394"])
x.add_row(["Linear SVM:OVR + Linear SVC","TFIDF","C=1.0","0.21","0.003"
,"0.70","0.25","0.37"])
x.add_row(["Linear SVM:OVR + SGDClassifier-Tuned","TFIDF","a=10**(-6)",
"0.15","0.004","0.46","0.32","0.38"])
print(x)

```

Values Obtained after working with 50K Datapoints :-

```

+-----+-----+-----+-----+-----+-----+-----+-----+
+---+---+---+---+---+---+---+---+
|                                     | Vectorizer | Para
m   | Acc  | Hamming | Pre  | Re   | Mic_F1 |
+---+---+---+---+---+---+---+---+
+---+---+---+---+---+---+---+---+
| Log Regression:OVR + SGDClassifier(loss='log') | TFIDF      | a=0.
01  | 0.21 | 0.003  | 0.67 | 0.27 | 0.38  |
|      Log Regression:OVR + LogisticRegression   | TFIDF      | -
|      | 0.16 | 0.003  | 0.80 | 0.16 | 0.22  |
|      Log Regression:OVR + SGDClassifier- Tuned   | TFIDF      | a=10^
(-6) | 0.18 | 0.003  | 0.55 | 0.31 | 0.393 |
|      Log Regression:OVR + SGDClassifier(loss='log') | BOW        | a=0.

```



01		0.06		0.009		0.16		0.36		0.22		
		Log Regression:OVR + LogisticRegression							BOW		C=1.	
0		0.17		0.003		0.52		0.28		0.360		
		Log Regression:OVR + LogisticRegression - Tuned							BOW		C=10	
0		0.14		0.004		0.45		0.30		0.361		
		Linear SVM:OVR + SGDClassifier(loss='hinge')							TFIDF		a=0.	
01		0.21		0.003		0.68		0.28		0.394		
		Linear SVM:OVR + Linear SVC							TFIDF		C=1.	
0		0.21		0.003		0.70		0.25		0.37		
		Linear SVM:OVR + SGDClassifier-Tuned							TFIDF		a=10**	
(-6)		0.15		0.004		0.46		0.32		0.38		
+-----+-----+-----+-----+-----+												
-----+-----+-----+-----+-----+												